

EINDHOVEN UNIVERSITY OF TECHNOLOGY

Reward Shape and Frame Stacking in DQN

A Reward $\times$ Observation Ablation in
Multi-Objective Grid Survival

BACHELOR'S THESIS

Philipp Merz

Supervisor: Maxence Delorme
Tilburg University

2026

Abstract

In this thesis we study two design choices that strongly affect how a deep reinforcement learning agent learns: the reward function it is asked to optimize, and how much short-term history of the world it observes at each step. The first determines what counts as good behavior. The second determines whether the agent can perceive motion at all, since a single snapshot only reveals position, but not direction or speed. To examine that interaction we built a small grid-world survival game in which an agent must eat to avoid starvation, evade enemies, and optionally take refuge in shelters, and trained a Deep Q-Network agent on this task while varying the reward function and the observation (either the current frame, or the four most recent frames concatenated). Our results show that frame stacking is not universally beneficial. It improved survival when maximizing the reward required actually surviving, but degraded it under two reward designs that admit easier ways to accumulate reward without surviving. Under those misspecified rewards, the agent used the additional perceptual capability to collect more reward at the cost of staying alive. The agent that combined the well-specified reward with the four-frame observation reached a survival level comparable to a scripted reference policy hand-written for the task. These findings are specific to our environment, but suggest that observation-side and reward-side design decisions cannot be treated as independent.

Contents

1	Introduction	2
2	Related Work	3
3	Background	4
3.1	Reinforcement Learning Fundamentals	4
3.2	Deep Q-Networks	5
3.3	DQN Extensions	5
3.4	Training Stability Techniques	6
3.5	Reward Shaping	7
4	Environment	7
4.1	World Design	8
4.2	Observation Space	9
4.3	Action Space	10
4.4	Reward Structure	10
5	Agent Architecture	11
5.1	Network	11
5.2	Training Algorithm	12
5.3	Hyperparameters	12
6	Experiments	13
6.1	Experimental Setup	13
6.2	Failure Modes	13
6.3	Reward Configurations	14
6.4	Weak Proximity: a Single-Seed Illustration	16
6.5	Reward $\times$ Observation Ablation	16
7	Results and Discussion	17
7.1	The three predictions	17
7.2	Behavior at evaluation	20
7.3	Learning dynamics	23
7.4	Limitations	23
8	Conclusion	25
A	Hyperparameters	27
B	Per-seed survival	28
C	Episode outcomes and shelter occupancy	28

1 Introduction

In reinforcement learning, an agent learns a policy (a mapping from observations of an environment to actions) by interacting with that environment. At each step it observes the current state, takes the action prescribed by its current policy, and receives a reward signaling how good the choice was. The agent’s objective is to discover a policy that, over time, accumulates as much reward as possible. The reward function therefore defines what the agent considers good behavior. When the task is simple, a single reward signal at the end of each episode can suffice. When the task is complex and combines several competing objectives (making progress, avoiding harm, managing limited resources), a single end-of-episode signal is too coarse, since the agent has no way to tell which of the many low-level actions taken during the episode contributed to the outcome.

The standard response is reward shaping, where the agent receives a small reward at every step in addition to or instead of the terminal signal, computed as a weighted sum of several quantities that proxy for the designer’s intent. Make progress toward the goal, gain reward. Take damage, lose reward. Stay near a safe area, gain reward. Each weighted component pulls the agent’s behavior in some direction, and collectively they are meant to steer it toward the intended overall policy. In practice, finding weights that produce the intended behavior is one of the hardest parts of training agents for multi-objective embodied control tasks such as legged-robot locomotion and aerial navigation [Hwangbo et al., 2019, 2017].

Small changes in relative weights can lead the agent to qualitatively different behaviors, including policies that maximize the shaped reward without producing what the designer had in mind. A reward designed to encourage the agent to forage may instead teach it to hover near food without eating it, because hovering pays continuously while eating ends the reward stream. This kind of mismatch between the shaped reward and the behavior it was meant to encode is known as reward hacking [Amodei et al., 2016]. From the agent’s perspective, no error has been made. It is simply optimizing the function it was given.

A second design choice, parallel to the reward, concerns what the agent perceives at each step. An environment that exposes only the current state allows the agent to access positional information, but not motion. Whether an approaching object is moving toward or away from the agent, for instance, cannot be inferred from a single snapshot. The standard remedy is to give the agent a short history of recent observations rather than only the most recent one, an approach known as frame stacking [Mnih et al., 2013]. With recent history available, the agent can in principle learn to infer derived quantities, such as velocity, directly from its input.

These two design axes, reward and observation, are routinely tuned together when training agents but are rarely studied as interacting variables in their own right. We approached the experiments described below with the expectation that providing more observation history via frame stacking should at worst do nothing, since extra information that the agent does not need can simply be ignored. We therefore hypothesized that the combination of the cleanest available reward design and the richer observation would produce the strongest agent, and that adding observation history would at worst be neutral under the other reward designs. The experiments will show that this expectation held only in part, and that the effect of frame stacking on survival depends on the shape of the reward in a surprising way.

To investigate this question we constructed a small grid-world survival environment in which an agent must manage hunger, evade enemies, and optionally take refuge

in shelters. The environment is deliberately simple, exposing only the mechanisms relevant to the reward-shaping question (competing objectives, a per-step managing of hunger and progress, and event-driven rewards for eating and being hit) and removing everything not relevant (rich action spaces, procedurally generated worlds, multi-agent interactions). On this environment we train one fixed agent architecture under several reward configurations and two observation depths, and compare its behavior against that of a hand-written reference policy that captures the intended use of the environment. The source code and experiment configurations are available at <https://github.com/philippmerz/rl-thesis>. The investigation is constrained to a single environment and a modest training budget, and the conclusions we draw are correspondingly modest. The detailed experimental design and methodology are given in Section 6.

The remainder of the thesis is organized as follows. Section 2 surveys the related literature. Section 3 reviews the technical background on reinforcement learning, deep Q-learning, and reward shaping. Section 4 describes the survival environment in detail and explains why we built a custom environment rather than reuse an existing benchmark. Section 5 details the agent we trained. Section 6 presents the experimental design, failure-mode analysis, reward configurations, and the main ablation. Section 7 reports the results and discusses them. Section 8 concludes.

2 Related Work

In contemporary deep reinforcement learning (RL) practice, reward shaping refers broadly to constructing a per-step reward as a weighted sum of components that each proxy for some aspect of the designer’s intent. Published reward functions for multi-objective embodied control commonly assemble several such components. The reward functions of Hwangbo et al. [2019] for legged-robot locomotion and Hwangbo et al. [2017] for quadrotor flight weigh forward progress, attitude stability, energy expenditure, smoothness, and other quantities against each other. Loquercio et al. [2021] use a comparable hand-tuned, multi-term objective for agile drone navigation. Across these examples the weights are hand-tuned, and their calibration is widely reported as one of the practical bottlenecks in training such policies. This thesis studies what happens when this reward-tuning task is paired with a second design decision on the observation side.

Amodei et al. [2016] catalogued reward hacking as one of several concrete safety failure modes that emerge when learned policies optimize a proxy reward instead of the underlying objective directly. The phenomenon predates deep RL by decades. Lehman et al. [2020] collect a wide range of earlier examples from evolutionary computation and artificial life, including virtual creatures evolved to jump that instead learned to exploit bugs in the physics simulation to gain height without ever actually jumping. Whether such exploits are found at all depends on the agent’s capability. Pan et al. [2022] make this dependence empirically concrete. Holding the reward misspecification fixed while varying model size, action-space resolution, observation noise, and training time, they find that more capable agents do better on the misspecified reward they were trained against but worse on the objective the reward was meant to capture, sometimes crossing thresholds at which the behavior changes abruptly. We later use this reward-capability interaction framework to interpret our own experimental findings.

Beyond the reward and the agent’s capability, the environment in which the agent is trained and tested is a third design lever. It must expose the phenomena under investigation in a form that can be measured cleanly. So-called gridworlds, simulated

2D-grid environments, have long served this purpose, since they isolate phenomena that are hard to study in more realistic and complex settings. Leike et al. [2017] introduce a suite of safety-focused gridworlds that expose common ways RL agents misbehave. Their environments include settings where the agent finds unintended shortcuts to high reward, settings where it causes collateral damage in pursuit of its objective, and settings where it performs badly when the deployment environment differs from the one it trained in. Several are explicit reward-hacking benchmarks of the kind described above. Hafner [2022] provide the Crafter benchmark, a multi-objective survival gridworld with a procedurally generated layout each episode in which the agent forages for food and water, sleeps, fights enemies, and builds shelter. The MiniGrid framework of Chevalier-Boisvert et al. [2023] is goal-oriented rather than survival-based, tasking the agent with reaching a target tile or retrieving an object across a modular set of layouts. The NetHack Learning Environment [Küttler et al., 2020] extends the survival paradigm to a much larger action space and far longer episodes. The environment used in this thesis is a leaner instance of this survival design.

Alongside the choice of environment is the choice of what the agent perceives within it. Frame stacking, the observation-side variable of our experiments, was introduced in the original deep Q-learning work [Mnih et al., 2013] to give a standard feed-forward network access to short-term motion information that a single frame cannot convey. The canonical alternative is to use recurrent network architecture instead. Hausknecht and Stone [2015] introduce DRQN, which extends DQN with such a recurrent component and remains the standard reference for the recurrent approach.

All of these design choices ultimately have to be exercised at some compute scale. The single-laptop training budget used in this thesis is supported by recent work showing that algorithmic choices can substitute for large-scale compute in deep Q-learning. Obando-Ceron and Castro [2021] revisit the deep Q-learning baseline of Hessel et al. [2018] under more careful evaluation. Schwarzer et al. [2023], Obando-Ceron et al. [2023], and Clark et al. [2025] demonstrate that consumer-grade hardware can match results previously requiring large-scale compute, through algorithmic refinements such as smaller networks and more efficient use of replay. These results inform the choice of agent and training budget in Section 5.

3 Background

3.1 Reinforcement Learning Fundamentals

We formulate the task as a Markov Decision Process (MDP) defined by a tuple (S, A, P, R, γ) : state space S , action space A , transition function $P(s' | s, a)$, reward function $R(s, a, s')$, and discount factor $\gamma \in [0, 1)$ [Sutton and Barto, 2018].

The agent interacts with the environment in discrete time steps $t = 0, 1, 2, \dots$, receiving reward $r_t \in \mathbb{R}$ at each step. The goal is to find a policy $\pi : S \rightarrow A$ that maximizes the expected discounted return $G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$. The action-value function

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t | s_t = s, a_t = a] \quad (1)$$

gives the expected return of taking action a in state s and following π thereafter. The optimal action-value function $Q^*(s, a) = \max_\pi Q^\pi(s, a)$ satisfies the Bellman optimality equation:

$$Q^*(s, a) = \mathbb{E} \left[r_t + \gamma \max_{a' \in A} Q^*(s_{t+1}, a') \right], \quad (2)$$

where s_{t+1} is the successor state reached after taking a in s .

Q-learning estimates Q^* directly through temporal difference (TD) updates:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_t + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right], \quad (3)$$

where α is the learning rate. The bracketed quantity is the TD error δ_t . These updates assume we can store and update a Q -value for every state-action pair, which is feasible only when the state space is small enough to enumerate.

3.2 Deep Q-Networks

Deep Q-Networks (DQN) [Mnih et al., 2013, 2015] address this limitation by approximating Q^* with a neural network $Q(s, a; \theta)$ parameterized by weights θ .

Two mechanisms stabilize training. First, instead of sampling temporally correlated consecutive state transitions, which would violate the optimizer’s i.i.d. data prerequisite, *experience replay* stores transitions (s, a, r, s') in a buffer and samples mini-batches uniformly. Second, a *target network* with parameters θ^- provides stable TD targets. The loss for a sampled transition is:

$$L(\theta) = \left[r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right]^2. \quad (4)$$

The target parameters θ^- are periodically copied from θ .

Exploration uses an ε -greedy policy: with probability ε a random action is taken, otherwise the greedy action $\arg \max_a Q(s, a; \theta)$.

Vanilla DQN as described above has known weaknesses, and a number of extensions have been proposed to address them individually. The next subsection describes the four such extensions that our agent incorporates.

3.3 DQN Extensions

The first weakness is that DQN systematically overestimates action values. The target uses $\max_{a'} Q(s', a'; \theta^-)$, so any positive estimation noise in Q is preferentially picked up by the max operator and biases the target upward. Double DQN [van Hasselt et al., 2016] decouples action selection from action evaluation:

$$y = r + \gamma Q\left(s', \arg \max_{a'} Q(s', a'; \theta); \theta^-\right). \quad (5)$$

The online network selects the maximizing action, and the target network evaluates it.

A second concern is that the network has to learn an explicit Q -value for every state-action pair, even in states where the action choice barely affects the outcome. The Dueling architecture [Wang et al., 2016] factors the Q -function into a state value $V(s)$ and an action advantage $A(s, a)$:

$$Q(s, a) = V(s) + A(s, a) - \frac{1}{|A|} \sum_{a'} A(s, a'), \quad (6)$$

where $|A|$ is the number of actions. The decomposition lets the network learn $V(s)$ without separately evaluating each action, which helps in states where the action choice has little effect on the value.

The third extension concerns how transitions are sampled from the replay buffer. Uniform sampling treats every transition the same, but transitions with large TD error carry more learning signal per sample. Prioritized experience replay (PER) [Schaul et al., 2016] samples transitions in proportion to their TD-error magnitude, revisiting surprising transitions more often. The sampling probability of transition i is

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}, \quad p_i = |\delta_i| + \epsilon_p, \quad (7)$$

where α controls prioritization strength and ϵ_p is a small constant ensuring nonzero probability. Non-uniform sampling introduces bias, which is corrected by importance-sampling (IS) weights

$$w_i = \left(\frac{1}{N \cdot P(i)} \right)^\beta \bigg/ \max_j w_j, \quad (8)$$

where N is the number of stored transitions and β is annealed from a starting value to 1 during training, gradually applying the full bias correction.

A fourth extension addresses the slow propagation of reward through single-step bootstrapping. n -step returns [Sutton, 1988, Sutton and Barto, 2018] accumulate rewards over n consecutive steps before bootstrapping:

$$R_t^{(n)} = \sum_{k=0}^{n-1} \gamma^k r_{t+k+1} + \gamma^n \max_{a'} Q(s_{t+n}, a'; \theta^-). \quad (9)$$

This propagates reward information faster at the cost of increased variance. Moderate values of n (typically 2–5) improve learning speed in practice.

Hessel et al. [2018] combined six such extensions into a single agent under the name Rainbow. The six are Double DQN, Dueling, PER, multi-step returns, distributional RL (C51), and NoisyNet. Their ablation showed each component contributes, with PER and multi-step returns providing the largest individual gains. We adopt the four techniques outlined above, but omit distributional RL and NoisyNet because of implementation complexity.

3.4 Training Stability Techniques

Beyond the algorithmic extensions above, two further design choices stabilize the optimization itself.

The original DQN copies θ to θ^- every C steps, which creates discontinuities in the target. The soft target update [Lillicrap et al., 2016] replaces the hard copy with an exponential moving average, updating the target at every step:

$$\theta^- \leftarrow (1 - \tau) \theta^- + \tau \theta, \quad (10)$$

with a small τ . This produces smoother target evolution rather than the periodic discontinuities of the hard copy.

Batch normalization, the standard choice in supervised deep learning, is unreliable in deep RL because the data distribution is non-stationary and effective batch sizes are small. LayerNorm [Ba et al., 2016] instead normalizes across features within each sample, independently of the batch. Gallici et al. [2024] show empirically across Atari-game ablations that LayerNorm improves deep Q -learning performance over batch normalization in this regime.

3.5 Reward Shaping

Reward shaping refers to the practice of designing a per-step reward signal that decomposes the intended overall objective into a sum of smaller, denser terms. When the task has competing sub-objectives, or when the terminal reward is too sparse for the agent to learn from in reasonable time, a hand-tuned dense reward gives step-by-step feedback that proxies the designer’s intent. The weights on the individual components are design parameters whose calibration determines whether the learned policy matches that intent.

Formally, a shaped reward is just another scalar function of state, action, and successor state. The agent’s only objective is to maximize its expected discounted sum. Whether that sum coincides with the designer’s intent depends entirely on whether the chosen function does. Most ad-hoc shaping is a weighted combination of heuristic components, each separately encoding a fragment of good behavior. The agent does not see those fragments as parts of an aggregate goal. It sees only a scalar at each step and optimizes the total. Any combination of weights that admits a higher-scoring strategy than the intended one creates a target for the agent to find, but whether the agent actually finds it depends on its capability.

Pan et al. [2022] make this dependence concrete. Holding misspecified rewards fixed while varying agent capability (model capacity, action-space resolution, observation noise, and training time), they find that more capable agents systematically earn higher *proxy* reward and lower *true* reward, in some cases crossing capability thresholds at which behavior changes abruptly. Reward hacking, in this view, is not a property of the reward function alone but a property of the function together with the agent’s capability to reach its loopholes. This makes for an unintuitive phenomenon: the same reward can produce intended behavior under a weak agent and pathological behavior under a stronger one.

Throughout the rest of this thesis, we use the term *leaky* informally to describe a reward function whose episode return can be inflated by behavior that does not serve the intended objective. An imperfect reward *leaks* return to behaviors outside what it was meant to incentivize, the way a leaky pipe loses water to places it was not meant to flow. The cleaner the reward (in the sense of admitting fewer such loopholes), the smaller the gap between proxy reward and true objective should be, and the less reward hacking should be possible.

4 Environment

The survival environment used throughout this thesis was built specifically for the work, implemented from scratch rather than as a fork or extension of an existing library. We chose this route because the closest available alternatives did not match the research question cleanly. The Crafter benchmark [Hafner, 2022] expands the survival theme into a much larger 17-action task with crafting, inventory, and 22 achievements, most of which are tangential to the reward-shape question studied here. MiniGrid [Chevalier-Boisvert et al., 2023] is narrower, but its built-in environments are explicitly goal-oriented, framed around reaching a target tile or picking up an object rather than indefinite survival under competing pressures. A custom implementation also lets the environment step fast enough that each of the twenty-four training runs in our ablation completes within a few hours on a consumer GPU. Figure 1 shows a representative mid-episode snapshot. The subsections below describe the design in detail.




Figure 1: Screenshot of the survival environment. The agent (blue dot) is near the grid center. Two enemies (red dots) are visible to the east. Food items (small green dots) are scattered across the grid and replenish stochastically over time. Shelters (large gray dots) are placed once at episode start and remain fixed. The bottom bar reports current health and hunger, and current tick count.

4.1 World Design

The environment is a 64×64 bounded grid with discrete time steps. An episode ends when the agent dies (health reaches zero) or when an episode cap on the number of ticks is reached. The world contains four entity types.

Agent. A single agent with two resource bars: health h and hunger u , both with maximum value $h_{\max} = u_{\max} = 100$. The agent spawns at the grid center with both bars full. Hunger depletes by 0.2 per tick and by an additional 0.5 when moving. When hunger reaches zero, the agent starves and loses 0.5 health per tick. When hunger exceeds 30% of its maximum (hunger ratio above $u_{\text{thresh}} = 0.3$), health regenerates at 0.5 per tick up to the cap.

Enemies. Mobile entities that deal 15 damage on contact (Manhattan distance ≤ 1). Each tick, an enemy moves with probability 0.5. Within a vision range of 3 (Manhattan distance), enemies pursue the agent along the axis of greatest distance. Outside this range, they move randomly. Enemies cannot enter shelter tiles or tiles occupied by another enemy. The world supports up to 8 enemies simultaneously, with roughly a third of this capacity present at episode start. New enemies appear with probability 0.01 per tick.

Food. Stationary items that restore 20 hunger when consumed by stepping onto their tile. Up to 36 food items exist simultaneously, with half present at episode start. New food appears with probability 0.05 per tick.

Shelters. Stationary tiles generated at episode start with density 0.01 per cell (approximately 41 shelters per episode). While in a shelter, the agent takes no damage from enemies. Shelters do not protect against starvation.

Each episode uses a different random seed, producing a unique shelter layout and initial entity placement. Enemy and food spawning continues throughout the episode, maintaining population pressure.

4.2 Observation Space

The agent observes through a 15×15 window (radius $r_{\text{obs}} = 7$) centered on the agent. The observation consists of three binary spatial channels:

1. Enemy presence
2. Food presence
3. Shelter presence

Cells outside the grid boundary are zero. Three scalar features are appended: $\frac{h}{h_{\text{max}}}$, $\frac{u}{u_{\text{max}}}$, and a binary shelter indicator. The full single-frame observation is a 678-dimensional vector ($3 \times 15 \times 15 = 675$ spatial entries plus 3 scalar features). The 4-frame stacked variant introduced as the observation-side ablation in Section 6.5 concatenates the spatial channels of the four most recent frames, raising the spatial channel count from 3 to 12. The spatial component therefore grows by an exact factor of 4 (from 675 to 2,700), but the three scalar features are taken from the current frame only, since they reflect the cumulative state of the agent and do not benefit from temporal stacking. Thus, the total observation dimension grows from 678 to 2,703, by a factor of $2703/678 \approx 3.99$. Figure 2 illustrates the layout.




Figure 2: Observation structure. The agent (blue dot at the window center) sees three binary 15×15 grids corresponding to enemy, food, and shelter presence within radius 7, plus three scalar features summarizing the agent’s internal state. Section 6.5 crosses this single-frame observation with a 4-frame stacked variant (12 spatial channels) as the observation-side ablation.

4.3 Action Space

The agent has five discrete actions: stay in place, or move one cell up, down, left, or right. Movement is clamped to the grid boundaries.

4.4 Reward Structure

The reward at each step is a weighted sum of up to eleven components:

$$\begin{aligned}
r_t = & w_{\text{food}} i_t^{(\text{eat})} + w_{\text{edm}} d_t^{(\text{e})} + w_{\text{sdm}} d_t^{(\text{s})} + w_{\text{death}} i_t^{(\text{die})} \\
& + w_{\text{hunger}} i_t^{(\text{low})} + w_{\text{hprop}} (1 - u_t) + w_{\text{tick}} \\
& + w_{\text{fprox}} \Delta \phi_t^{(\text{f})} + w_{\text{eprox}} \Delta \phi_t^{(\text{e})} + w_{\text{sprox}} \Delta \phi_t^{(\text{s})} \\
& + w_{\text{shelter}} i_t^{(\text{safe})},
\end{aligned} \tag{11}$$

where $w. \in \mathbb{R}$ are tunable scalar weights; $i_t^{(\text{eat})}, i_t^{(\text{die})}, i_t^{(\text{low})}, i_t^{(\text{safe})} \in \{0, 1\}$ are binary indicators for food eaten, death, low hunger, and shelter safety; $d_t^{(\text{e})}$ and $d_t^{(\text{s})}$ are enemy damage and starvation damage; and $u_t = u/u_{\text{max}} \in [0, 1]$ is the hunger ratio.

Table 1 lists all reward components and their baseline weights.

Table 1: Reward components and baseline weights. Components marked “disabled” are set to zero in the baseline and activated in specific experimental configurations.

Component	Condition	Baseline
Food eaten (w_{food})	Agent steps on food tile	+5.0
Enemy damage (w_{edm})	Per unit of damage received	−0.1
Starvation damage (w_{sdm})	Per unit of starvation damage	−1.0
Death (w_{death})	Health reaches zero	−10.0
Low hunger (w_{hunger})	Per tick when $u_t < u_{\text{thresh}}$	disabled
Hunger proportional (w_{hprop})	Continuous, scales with $(1 - u_t)$	−0.3
Food proximity (w_{fprox})	Change in closeness to nearest food this step	+0.1
Enemy proximity (w_{eprox})	Change in closeness to nearest enemy this step	disabled
Shelter proximity (w_{sprox})	Change in closeness to nearest shelter this step	disabled
Shelter safety (w_{shelter})	Agent in shelter, enemy within distance 3	+0.1
Survival tick (w_{tick})	Per step alive	disabled

The three proximity terms reward per-step changes in closeness to a target, rather than absolute closeness. For target type $x \in \{\text{f}, \text{e}, \text{s}\}$ (food, enemy, shelter), define the closeness score

$$\phi_t^{(x)} = \frac{\text{dist}_{\text{max}} - \text{dist}_t^{(x)} + 1}{\text{dist}_{\text{max}} + 1} \in [0, 1], \tag{12}$$

where $\text{dist}_t^{(x)}$ is the Manhattan distance to the nearest visible entity of type x and $\text{dist}_{\text{max}} = 2r_{\text{obs}}$ (maximum Manhattan distance to any cell within the observation radius). If no entity is visible, $\phi_t^{(x)} = 0$. The per-step closeness change $\Delta \phi_t^{(x)} = \phi_t^{(x)} - \phi_{t-1}^{(x)}$ is positive when the agent moved closer to the target in that step and negative when it moved away. We reward this per-step change rather than the absolute closeness because rewarding absolute closeness would also reward standing still next to a target (for example, hovering next to a food tile without eating it). The per-step change form gives zero reward for staying in place, so it only pays the agent for moving toward the target.

Two proximity terms can be conditionally turned on (gated) based on the hunger ratio, using a dedicated gating threshold u_{gate} . When conditional gating is enabled, the food proximity fires only when $u_t < u_{\text{gate}}$ (hungry) and the shelter proximity fires only when $u_t \geq u_{\text{gate}}$ (well-fed). This encodes a behavioral mode switch through the reward function: the agent receives a spatial indicator toward food when it needs to eat and toward shelter when it does not.

These components create competing pressures. Collecting food requires movement, which costs hunger and exposes the agent to enemies. Shelters neutralize enemy damage but provide no food. The hunger-proportional penalty grows every tick, creating urgency to eat that the agent cannot escape through passivity.

Another design consideration is the relative magnitude of continuous per-tick rewards versus sparse event rewards. A small per-tick reward that fires at every step can in aggregate exceed an infrequent event reward, biasing the agent toward passive survival rather than active foraging. Similarly, a food proximity reward that is too large relative to the food-eaten reward incentivizes hovering near food without eating it, since eating removes the food and terminates the proximity stream. These failure modes motivate the experimental configurations in Section 6.

These weights are the primary independent variable in our experiments.

5 Agent Architecture

The agent combines the four DQN extensions and two stability techniques introduced in Section 3 (Double DQN, the Dueling architecture, prioritized experience replay, n -step returns, Polyak target updates, and LayerNorm) with AdamW [Loshchilov and Hutter, 2019] as the optimizer. This combination is sometimes called Rainbow-lite, a stripped-down Rainbow that omits more elaborate components for implementation simplicity. The subsections below describe the network architecture, the training procedure, and the hyperparameters used. Full per-parameter values are in Appendix A.

5.1 Network

The agent is implemented in PyTorch [Paszke et al., 2019]. Because this thesis varies reward shape and observation space rather than network architecture, the architecture must be adequate but does not need to be optimal. The key requirement is enough representational capacity that the network is not the bottleneck under any of the configurations we test.

A convolutional neural network is the natural choice because the observation has spatial structure ($3k \times 15 \times 15$ grid for frame-stack size k). The architecture follows the Atari-DQN backbone of Mnih et al. [2015], retained by subsequent DQN-extension papers [Wang et al., 2016, Hessel et al., 2018], and consists of three convolutional layers followed by a fully connected layer projection and the Dueling value-advantage head from Section 3.3. Internal dimensions are scaled down from the Atari version to fit the smaller 678-dimensional flat observation, and were not separately tuned for this task. The comparisons reported below depend on the architecture being held constant across cells, not on it being optimal.

For single-frame input ($k = 1$), the network has 970,438 trainable parameters mapping the 678-dimensional observation to 5 output actions. The only architectural change in the frame-stacked cells ($k = 4$) is the input depth of the first convolutional layer. Full layer specifications are in Appendix A.

5.2 Training Algorithm

Algorithm 1 summarizes the training procedure.

Algorithm 1 Rainbow-lite DQN training.

- 1: Initialize policy network θ , target network $\theta^- \leftarrow \theta$
 - 2: Initialize prioritized replay buffer $\mathcal{D}$, n -step accumulator
 - 3: Collect 10,000 transitions with random policy into $\mathcal{D}$
 - 4: **for** $t = 1, \dots, T$ **do**
 - 5: Select a_t via ε -greedy from $Q(\cdot; \theta)$, decay ε
 - 6: Execute a_t , observe r_t, s_{t+1}, d_t
 - 7: Compute n -step return; commit to $\mathcal{D}$ when ready
 - 8: Sample batch from $\mathcal{D}$ with PER; compute importance-sampling (IS) weights w_i
 - 9: $y_i = R_i^{(n)} + (1 - d_i) \gamma^n Q(s'_i, \arg \max_{a'} Q(s'_i, a'; \theta); \theta^-)$
 - 10: $L = \frac{1}{B} \sum_i w_i \cdot \text{Huber}(Q(s_i, a_i; \theta), y_i)$
 - 11: Update θ via AdamW optimizer
 - 12: $\theta^- \leftarrow (1 - \tau) \theta^- + \tau \theta$
 - 13: Update priorities in $\mathcal{D}$ from new TD errors
 - 14: **end for**
-

Here, $d_i \in \{0, 1\}$ indicates whether sampled transition i is terminal, B denotes the batch size, and Huber is the Huber loss, which uses squared error for small residuals and absolute error for large ones, reducing sensitivity to outliers in the TD target compared to mean squared error.

At each step, the transition $(s_t, a_t, r_t, s_{t+1}, d_t)$ enters an n -step accumulator. Once $n = 3$ transitions have accumulated, the oldest is committed to the replay buffer with return $R_t^{(3)} = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2}$, bootstrap state s_{t+3} , and discount factor γ^3 . At episode termination (agent death), remaining entries are flushed with shortened returns. On episode truncation (time limit), pending entries are discarded.

The distinction between termination and truncation matters for the bootstrap target. A terminated episode (death) sets $d = 1$, zeroing the bootstrap term. A truncated episode sets $d = 0$, allowing bootstrapping from the final state since the agent was still alive.

The replay buffer stores transitions in pre-allocated NumPy arrays [Harris et al., 2020]. To support the priority-proportional sampling required by PER, it uses a SumTree data structure [Schaul et al., 2016], a binary tree that stores partial sums of priorities at internal nodes and supports both $O(\log n)$ priority updates and $O(\log n)$ proportional sampling. Stratified sampling further divides the cumulative priority range $[0, P_{\text{total}}]$ into equal segments and draws one sample per segment, reducing variance compared to pure proportional sampling.

The optimizer is the standard AdamW [Loshchilov and Hutter, 2019], with the learning rate held constant.

Exploration follows an ε -greedy schedule with linear decay from 1.0 to 0.01 over the first 500,000 steps (one quarter of training), then held constant. The non-zero floor of 0.01 preserves a small amount of exploration throughout the rest of training.

5.3 Hyperparameters

Appendix A lists all hyperparameters. Values follow established defaults from the literature rather than environment-specific tuning. The frame-stack size k is the observation-side

ablation variable (Section 6.5). Every other hyperparameter is held constant across all twenty-four ablation runs.

6 Experiments

6.1 Experimental Setup

Each ablation run trains for $T = 2,000,000$ timesteps (a few hours per seed on a consumer GPU). The misspecified configurations (BASELINE, ABSOLUTE_PROXIMITY, and their frame-stacked variants) cap episodes at 1,000 ticks, whereas the MINIMAL configurations cap at 50,000 ticks. An episode cap is required by the protocol, since a policy could in principle survive indefinitely, which would prevent benchmark and training termination. We initially used a 1,000-tick cap throughout, chosen to leave roughly 30% headroom above the heuristic’s mean survival of ≈ 769 ticks. Preliminary results showed that the misspecified configurations reached the cap in only a negligible fraction of episodes, because their reward-hacking failure modes terminate episodes well before the cap. We accordingly left their cap at 1,000. The MINIMAL configurations, in contrast, reached the cap in a non-negligible fraction of episodes (the minimal reward is not intentionally misspecified, so the agent survives long enough for right-censoring at the lower cap to materially affect the cell mean), so we raised their cap to 50,000 ticks. The heuristic baseline is benchmarked at both caps so that each cell is paired with a heuristic measurement at the same cap. The heuristic’s mean survival is essentially unchanged between the two caps (768.8 at cap 1,000 vs 772.1 at cap 50,000). Greedy ($\varepsilon = 0$) in-training evaluation runs every 5,000 steps with 100 episodes. Final benchmarks evaluate the best in-training checkpoint over 100 episodes with fixed seeds (1000–1099). We report mean survival ticks (with 95% confidence interval), mean food eaten, mean damage taken, and death rate. Each cell is run with four random seeds (42, 43, 44, 45). The agent architecture, all hyperparameters in Appendix A, and the training budget are held constant. The only variables are the reward configuration (Section 6.3) and the frame-stack size $k \in \{1, 4\}$ (Section 6.5).

As a non-learned reference we compare against a scripted heuristic agent. The heuristic camps in the nearest shelter and leaves only when the hunger ratio drops below 0.5, at which point it moves toward the nearest food. It flees from enemies within Manhattan distance 5. Over 100 evaluation episodes, this agent survives a mean of 769 ± 19 ticks (95% CI), eats 1.19 food per episode, and dies in 94% of episodes. The remaining 6% reach the 1,000-tick cap. The heuristic is rerun on the same evaluation seeds for every cell, so heuristic and DQN episodes are matched by evaluation world. Each cell mean carries a standard error taken across its four per-seed means, shown as the error bars of Figure 3. For each cell, we compute one number per world by averaging across its four training seeds, giving 100 per-world cell values. Subtracting the heuristic’s survival on each world gives 100 differences. We report the mean and 95% confidence interval of those differences as the primary statement, with a two-sided p -value alongside.

6.2 Failure Modes

Early experiments revealed two degenerate failure modes that emerged reliably under certain reward configurations. The two can co-occur within a single training run.

Exposed stasis. The agent fails to forage successfully and eventually dies of starvation in the open. The umbrella covers two recognizable sub-patterns in the action trace. The first is *oscillation*, where the agent cycles between adjacent cells, moving and undoing the move repeatedly (positions $A, B, A, B, \dots$). The second is *food hovering*, where the agent sits next to a food tile without crossing onto it. Hovering can arise under absolute closeness (ϕ rather than $\Delta\phi$) when the per-tick proximity reward exceeds the per-tick hunger penalty, in which case standing next to food earns reward indefinitely while eating it would terminate the stream. Avoiding this requires keeping the proximity weight low enough that the inequality flips, which is the design decision ABSOLUTE_PROXIMITY makes at $w_{\text{prox}} = 0.1$ (Section 6.3). In all variants of exposed stasis, the agent ends up dying of starvation outside a shelter.

Suicide. When the proximity reward is too weak to guide the agent toward food, the agent cannot learn to forage. Without foraging, every tick accumulates hunger-proportional penalty. A stationary agent depletes hunger at 0.2 per tick from a starting value of 100, reaching starvation at tick 500 and dying 200 ticks later, since health is lost at 0.5 per tick once starving. Over those 700 ticks, the undiscounted reward sums to:

- A hunger-proportional penalty of $\sum_{t=0}^{700} w_{\text{hprop}}(1 - u_t) \approx -135$. Over the first 500 ticks, the per-tick penalty rises linearly from 0 to -0.3 as hunger depletes from full to empty, accumulating to $\frac{1}{2} \cdot 500 \cdot -0.3 = -75$. The final 200 ticks at full starvation each contribute -0.3 , adding -60 .
- A starvation-damage penalty of $200 \cdot w_{\text{sdm}} \cdot -0.5 = -100$ for the 200 starvation ticks, each costing 0.5 health at reward weight $w_{\text{sdm}} = -1$.
- A death penalty of $w_{\text{death}} = -10$.

The total is approximately -245 . The alternative is to die quickly by seeking enemies, where each enemy hit deals 15 damage at $w_{\text{edm}} = -0.1$, costing -1.5 per hit. The agent reaches an enemy and dies in a few hits, accumulating roughly -10 to -15 in enemy damage plus the -10 death penalty, for a total cost of ≈ -24 . Slow starvation therefore costs roughly an order of magnitude more reward than seeking a quick death, which is the rational response to a reward function where living is more expensive than dying.

Both failure modes are clear instances of reward hacking [Amodei et al., 2016, Leike et al., 2017], in which the agent optimizes the reward function as specified but the resulting behavior diverges from the designer’s intent. They motivate the reward configurations introduced in Section 6.3. Section 7.2 quantifies both along two separate axes: every evaluation episode is classified by its cause of death, which distinguishes suicide from ordinary deaths to enemies and starvation, and the per-tick movement composition of each cell measures the stasis sub-patterns directly.

6.3 Reward Configurations

The failure modes above show that poorly chosen reward weights produce degenerate policies. Drawing on this analysis, we test four reward configurations: BASELINE as the reference, ABSOLUTE_PROXIMITY which alters the form of the proximity term, MINIMAL which strips the reward to three closeness-change proximity terms plus the death penalty, and WEAK_PROXIMITY as a single-seed illustration of what happens when the foraging signal is too weak. Table 2 summarizes them; each is motivated in turn below.

Table 2: Reward configurations used in the experiments. Only weights that differ from Table 1 are shown.

Configuration	Modification	Role
BASELINE	(none; default weights)	ablation cell, default reward
ABSOLUTE_PROXIMITY	$\Delta\phi$ replaced by ϕ (no closeness change)	ablation cell, tests the closeness-change reward shape
MINIMAL	minimal four-component reward (Table 3)	ablation cell, isolates proximity-change rewards
WEAK_PROXIMITY	$w_{\text{fprox}} = 0.02$	illustration of the suicide failure mode (single seed)

Baseline. The default weights (Table 1). The hunger-proportional penalty (-0.3) creates urgency, the food proximity reward ($+0.1$) provides a directional signal, and the food, damage, shelter, and death components encode the rest of the survival objective. These specific values were arrived at through preliminary experimentation and intuition during early development rather than through systematic tuning. The configuration is dense and compositional: every entity type contributes a reward signal and every internal state has a corresponding term. This makes BASELINE the most informative reward in principle, but also the most exposed to interactions between components.

Absolute proximity. The proximity terms are switched from per-step closeness change $\Delta\phi$ to absolute closeness ϕ . Every other weight matches BASELINE. This configuration tests the importance of the closeness-change reward shape. Under $\Delta\phi$, the cumulative proximity reward over a sequence of steps equals $w_{\text{fprox}} \cdot (\phi_{\text{final}} - \phi_{\text{initial}})$, bounded by w_{fprox} regardless of how many steps the agent takes: a round-trip back to the starting position earns zero, and the only way to earn proximity reward is to end the trajectory closer to a target than where it started. Under ϕ , a stationary agent at distance d from food earns $w_{\text{fprox}} \cdot \phi(d)$ per tick indefinitely, so cumulative proximity reward grows linearly with time spent near food, without bound. With the default $w_{\text{fprox}} = 0.1$ the per-tick proximity reward (≤ 0.1) still falls below the average hunger penalty (0.15), so hovering is not strictly net-positive per tick. But the incentive it creates is qualitatively different from BASELINE’s, and that difference is what this cell isolates.

Minimal. A minimal reward set with only three closeness-change proximity rewards and the death penalty:

- food-proximity ($w_{\text{fprox}} = +0.15$), gated to fire only when $u_t < 0.5$ (hungry),
- shelter-proximity ($w_{\text{sprox}} = +0.15$), gated to fire only when $u_t \geq 0.5$ (well-fed),
- enemy-proximity ($w_{\text{eproxy}} = -0.5$), always active,
- death ($w_{\text{death}} = -10$).

Every other weight is set to zero. The motivation is to cleanly separate the movement-based rewards from the event-driven rewards (food eaten, damage taken, shelter safety) and from the always-on hunger penalty. Each of the remaining three signals corresponds to a single objective, and the gating mirrors the heuristic’s mode switch (forage when hungry, shelter when not). Setting w_{hprop} and w_{sdm} to zero also removes the suicide pressure described in Section 6.2: with both terms zeroed, slow starvation and fast death cost the same terminal -10 , so the difference in accumulated penalty between the two strategies that incentivizes a randomly-initialized policy to die disappears. Table 3 lists the resulting weights.

Table 3: MINIMAL weights. All unlisted components from Equation 11 are set to zero. Food and shelter proximity are conditionally gated by $u_{\text{gate}} = 0.5$, configured independently from the health-regeneration threshold of Section 4.1 (which remains at its default value of 0.3).

Component	Condition	Weight
Food proximity (w_{fprox})	Closeness change, active only when $u_t < 0.5$	+0.15
Shelter proximity (w_{sprox})	Closeness change, active only when $u_t \geq 0.5$	+0.15
Enemy proximity (w_{eprox})	Closeness change, always active	−0.50
Death (w_{death})	Terminal	−10.0

Weak proximity (illustration). The rewards are the same as in BASELINE, but the food proximity weight is reduced from 0.1 to 0.02. At this weight, one step toward food earns approximately $0.02/15 \approx 0.00133$ in proximity reward (a single cell of approach is $1/15$ of the closeness range, from Section 4.4), which is smaller than the extra hunger penalty that move imposes on every future tick ($|w_{\text{hprop}}| \cdot 0.5/u_{\text{max}} = 0.0015$ per tick, from the 0.5 extra hunger drop each move costs). At random initialization, before the agent has learned that approaching food leads to eating it, the expected reward from moving toward food does not outweigh this hunger cost, and the agent prefers to die quickly rather than slowly starve, which is the suicide failure mode of Section 6.2. This is included as a single-seed illustration, not as part of the main ablation. Section 6.4 reports the resulting agent behavior qualitatively.

6.4 Weak Proximity: a Single-Seed Illustration

At random initialization, before the agent has learned that approaching food leads to eating it, the expected reward from moving toward food does not outweigh the extra hunger cost from movement (Section 6.3). The agent should therefore prefer to die quickly rather than slowly starve. We trained a single seed of WEAK_PROXIMITY ($w_{\text{fprox}} = 0.02$, all other weights matching BASELINE) for two million steps and benchmarked the result over 100 episodes. The agent reached mean survival 647 ± 33 ticks with 0.66 food per episode and a 99% death rate. These aggregate statistics match the four-seed BASELINE cell mean to the precision reported. The earlier argument assumed the agent picks the locally best action, but at random initialization it does not. ε -greedy starts at $\varepsilon = 1.0$, so every action is random until the network learns something. These random actions bump the agent into food often enough to discover the +5 eating reward by accident, after which it learns to forage regardless of the proximity reward’s size. We do not draw conclusions from this single seed beyond that observation.

6.5 Reward $\times$ Observation Ablation

The main experiment crosses three reward configurations (BASELINE, ABSOLUTE_PROXIMITY, MINIMAL) with two observation types (single-frame, $k = 1$, and 4-frame stacked, $k = 4$), four random seeds per cell. Each run trains for two million timesteps under the hyperparameters in Appendix A. The only varying parameters are the reward weights (Section 6.3) and k . The benchmark protocol is described in Section 6.1.

Three expectations motivated the experiments, labelled P1–P3:

- (P1) MINIMAL with frame stacking is the highest-performing of the six cells.
- (P2) MINIMAL with frame stacking exceeds the heuristic.




Figure 3: Reward × observation ablation. Each bar in the first two panels is the mean over four seeds. Black dots show individual seeds. Error bars are the standard error of the mean. Dashed lines mark the heuristic baseline (769 ticks survival, 1.19 food per episode). The third panel shows the mean change in survival between the single-frame and 4-frame configurations of each reward.

(P3) Frame stacking improves every reward configuration over its single-frame counterpart.

The reasoning behind P1 and P2 was that the cleanest reward (MINIMAL) combined with the temporal context required to infer enemy velocity (4-frame stacking) should produce the strongest policy in the experiment, and that this strongest policy would be sufficient to beat the scripted heuristic. P3 captured the prior expectation that adding observation history is at worst neutral for the aforementioned reasons. These three predictions are what the ablation was designed to test. As Section 7.1 reports, the experiment affirmed P1, left P2 unresolved, and refuted P3. It also revealed a structured effect that none of the three predictions anticipated.

7 Results and Discussion

We organize the following around the three predictions of Section 6.5, then turn to what the trained agents actually do at evaluation, the in-training learning dynamics, and the limitations of the study.

7.1 The three predictions

Table 4 reports per-cell averages across four seeds. Table 7 reports the per-seed numbers. Figure 3 summarizes survival, foraging activity, and the frame-stacking effect.

Prediction P1 held. Mean survival across the six cells ranges from 555 (BASELINE_FS) to 762 (MINIMAL_FS), and MINIMAL_FS is the only cell that comes close to the heuristic. The second-best cell is MINIMAL (single-frame) at 684 ticks, trailing MINIMAL_FS by 78 ticks. The within-cell standard error of MINIMAL_FS is small enough that no other cell is consistent with matching its mean.

P2 is not affirmed. Frame stacking lifts MINIMAL from 684 to 762 ticks, but the paired

Table 4: Reward $\times$ observation ablation. Each cell aggregates four seeds (42, 43, 44, 45). “Paired diff vs heur” is the cell mean (over training seeds) minus the heuristic, averaged across the 100 matched evaluation worlds, with the 95% confidence interval on that difference in brackets. WEAK_PROXIMITY is a single-seed reference and is reported separately. The MINIMAL cells (sf and fs) are evaluated at the 50,000-tick cap and paired with the heuristic at the same cap. All other cells are evaluated at the 1,000-tick cap. See Section 6.1 for the rationale.

Cell	Survival (mean)	Food	Death rate	Paired diff vs heur (95% CI)
Heuristic	769 $\pm$ 19	1.19	94%	—
<i>Single-frame ($k = 1$)</i>				
BASELINE	647	0.66	98%	−122 [−149, −94]
ABSOLUTE_PROXIMITY	634	0.43	100%	−135 [−166, −103]
MINIMAL	684	0.17	100%	−88 [−111, −64]
<i>4-frame stacked ($k = 4$)</i>				
BASELINE_FS	555	4.83	97%	−214 [−242, −187]
ABSOLUTE_PROXIMITY_FS	587	4.82	96%	−182 [−209, −155]
MINIMAL_FS	762	2.54	100%	−10 [−36, +16]
WEAK_PROXIMITY (1 seed)	647	0.66	99%	<i>n/a (single seed)</i>

gap to the heuristic is -10.0 ticks (95% CI $[-36, +16]$, $p = 0.45$). The confidence interval includes zero, so MINIMAL_FS could be tied with the heuristic, but we cannot show that it exceeds the heuristic. The original aim of the project was to raise DQN survival to or beyond the heuristic’s level, and the data are consistent with that aim being met within statistical noise rather than affirming a clear excess.

Before turning to P3, the spread of the single-frame row is itself worth noting. The three single-frame cells span only about 40 ticks of survival (634 for ABSOLUTE_PROXIMITY, 647 for BASELINE, 684 for MINIMAL), and all three fall below the heuristic at $p < 0.001$. Reward configuration matters marginally in this row, but not enough to bring DQN within reach of the heuristic without temporal context.

P3 did not hold, and the direction flipped. With $\Delta = \text{mean}(\text{fs}) - \text{mean}(\text{sf})$, we observe $\Delta = +78$ for MINIMAL, $\Delta = -47$ for ABSOLUTE_PROXIMITY, and $\Delta = -92$ for BASELINE (Figure 3, third panel). Frame stacking helps the cleanest reward and hurts the two leakier ones, both by margins much larger than within-cell variance. In principle, additional observation channels should weakly improve the agent, since extra information can be ignored if it is not useful. The data are inconsistent with this expectation, and this flip is the largest departure from our priors and motivates the rest of the discussion.

The obvious first explanation for the leaky-cell degradation is capacity. The observation grows $3.99\times$ while the network is held fixed, so the frame-stacked agent could simply be under-provisioned. But that explanation is reward-independent, and would thus predict uniform degradation across the three rewards, whereas the observed effect flips sign with reward shape and MINIMAL_FS is the strongest cell in the study. The effect of frame stacking is mediated by the reward, not the architecture.

While the flip cannot be explained by capacity, the episode return clarifies the mechanism. Within each reward configuration, the reward function is identical across k , so episode return is directly comparable across observation types. Frame stacking raises mean return on both leaky rewards (by about $+38$ on BASELINE and $+17$ on ABSOLUTE_PROXIMITY) while lowering survival. On MINIMAL the return is essentially unchanged across k (it sits near -10 regardless of k , since the minimal reward is closeness-change proximity rewards




Figure 4: Frame-stacking effect on mean episode return (x-axis) and mean survival ticks (y-axis), one point per reward configuration. Each point is the (sf $\rightarrow$ fs) shift in the per-cell means, averaged across the four training seeds and 100 evaluation worlds. The shaded lower-right region marks the reward-hacking quadrant where return rises and survival falls.

plus a terminal death penalty, neither of which scales with episode length) and survival rises. Figure 4 plots the frame-stacking delta in return against the delta in survival for each reward. The two leakier rewards land in the lower-right quadrant where return rises and survival falls, the geometric signature of the agent collecting more reward while staying alive less long. MINIMAL sits on the y-axis with a positive survival shift and essentially no return shift. Under the leaky rewards, the frame-stacked agent is more effective at collecting reward than the single-frame agent, but at the cost of survival. The reward function is misaligned with survival, and the agent is faithfully optimizing it. This is the pattern of capability-driven reward hacking [Pan et al., 2022].

The pattern we observe is one instance of the reward-capability interaction documented across four richer environments by Pan et al. [2022]. The frame-stack increase from $k = 1$ to $k = 4$ functions as a capability axis comparable to those they study (model capacity, action-space resolution, observation noise, and training time), and our two leaky reward configurations exhibit the same direction they document across all of theirs. More capability brings higher proxy reward and lower true reward, and frame stacking exposes the exploitable surplus in the leakier reward designs.

Pan et al. additionally identify *phase transitions*, “capability thresholds at which the agent’s behavior qualitatively shifts, leading to a sharp decrease in the true reward.” Our experimental design varies frame stacking only between $k = 1$ and $k = 4$ and is therefore insufficient to resolve whether the survival drop is smooth in k or has a phase-transition structure at some intermediate value. The broader implication for reward-design practice is that a reward function that looks adequate at one level of agent capability may produce a qualitatively different policy at another, particularly in applied settings where the true reward is observed only sporadically and the proxy is what is trained against.




Figure 5: (Left) episode outcomes of the 100-episode benchmark, averaged across the four seeds for each cell. Every episode reaches the time cap or ends in death from enemy contact or starvation. Enemy deaths split into enemy suicide (the agent’s own moves brought it closer to the enemy over the final encounter) and failure to escape (it retreated or stood still). (Right) mean shelter occupancy (fraction of survived ticks spent inside a shelter). The heuristic spends $\approx 82\%$ of its life in shelters. Every DQN cell spends $< 2\%$ regardless of reward configuration.

7.2 Behavior at evaluation

To characterize what each cell is doing beyond aggregate survival, we separate how episodes end from what the agent does with its time. Every episode either reaches the time cap or ends in death, and every death is attributed to its terminal damage, which is either enemy contact or starvation. Deaths to enemy contact are split further by the agent’s behavior over the *final encounter*, the unbroken run of ticks with an enemy in view that ends at death. Over that window we sum the per-tick change in distance to the nearest visible enemy caused by the agent’s own move, holding enemy positions fixed within each tick so that enemy pursuit is not misattributed to the agent. A negative sum means the agent’s own moves brought it closer to the enemy, and the episode counts as enemy suicide, the suicide mode of Section 6.2. A zero or positive sum means the agent stood still or retreated, and the episode counts as a failure to escape. This yields four mutually exclusive outcome categories (time cap, enemy suicide, failure to escape, starvation), shown in Figure 5, with exact percentages in Appendix C.

The second view is the movement composition of Table 5, the mean fraction of episode ticks the agent spends stationary, oscillating, or moving otherwise. An oscillation tick is one where the agent repeatedly moved and undid the move ($p_t = p_{t-2}$ with $p_t \neq p_{t-1}$).

Figure 6 shows one representative episode per cell, plotted as the agent’s position trace through the 64×64 grid. The top-row (single-frame) panels collapse to a small cluster of cells. The bottom-row (frame-stacked) panels both cover much more of the grid and show the resulting increase in food eaten.

Two regularities stand out in the displayed outcomes. First, starvation is the dominant death cause everywhere, from 55–66% of episodes in the leaky cells to 92–96% for the heuristic and the MINIMAL cells. Second, enemy suicide shows exactly what the analysis of Section 6.2 predicted: the two leaky rewards retain the hunger-proportional and starvation

Table 5: Movement composition. Mean percentage of episode ticks spent stationary, oscillating (sustained two-cell cycles, at least A, B, A, B), or in other movement, averaged over episodes and seeds.

Cell	Stationary	Oscillating	Other movement
Heuristic	98.1%	0.2%	1.8%
BASELINE	90.4%	6.2%	3.5%
ABSOLUTE_PROXIMITY	91.2%	5.9%	2.9%
MINIMAL	96.0%	1.8%	2.2%
BASELINE_FS	47.7%	7.7%	44.6%
ABSOLUTE_PROXIMITY_FS	51.1%	9.0%	39.9%
MINIMAL_FS	78.0%	2.3%	19.8%

penalties that make a quick death cheaper than slow starvation, and their cells show 19–25% enemy suicide, with every one of the sixteen leaky-cell seeds between 12% and 39%. MINIMAL zeroes those penalties, equalizing the reward cost of fast and slow death, and its suicide rate is 0% single-frame and 1.8% frame-stacked.

The movement composition captures what the outcomes barely register, namely the effect of frame stacking. Within each reward, the outcome mix changes only modestly from $k = 1$ to $k = 4$, but the movement mix transforms. The single-frame cells spend 90–96% of their ticks stationary, even more immobile than the heuristic (98%), which at least stands still inside shelters. With frame stacking, stationary ticks fall to 48–51% in the leaky cells and 78% in MINIMAL, and most of the freed time goes to genuine movement rather than oscillation. Oscillation is a minor pattern everywhere, at most 9% of ticks (ABSOLUTE_PROXIMITY_FS), so exposed stasis in this environment is overwhelmingly standing still rather than cycling in place. In short, the two axes act on different aspects of behavior. Outcomes follow the reward: the leaky cells differ from MINIMAL in their suicide and starvation shares at both observation depths, while going from $k = 1$ to $k = 4$ barely changes the outcome mix within a reward. Movement follows the observation: stationary time falls under frame stacking for every reward. The two axes still interact in the size of that fall: about 40 percentage points under the leaky rewards against 18 under MINIMAL, which leaves MINIMAL_FS far more stationary (78%) than the leaky frame-stacked cells (48–51%), consistent with its higher survival, since exposure scales with time spent moving in the open.

The added movement is productive foraging, not wandering. Food per episode rises from 0.17 to 2.54 for MINIMAL, from 0.43 to 4.82 for ABSOLUTE_PROXIMITY, and from 0.66 to 4.83 for BASELINE. With frame stacking, the agent learns to follow the food-proximity reward more aggressively, but in the leaky reward configurations the resulting policy spends more ticks exposed to enemies and dies sooner. Foraging behavior is acquired and safety behavior is not. Only MINIMAL_FS reaches a foraging rate consistent with sustained survival (2.54 food per episode at survival within 10 ticks of the heuristic, against the heuristic’s 1.19 food at 769 ticks).

Foraging is one half of the heuristic’s strategy. Shelter use is the other half, and it is essentially absent across all six DQN cells (median 0.9% of survived ticks, max 1.4% for ABSOLUTE_PROXIMITY_FS, driven largely by one outlier seed averaging 3.0% over its 100 episodes while the other three sat at 0.5–1.6%). No cell has learned to use the shelters as the heuristic does at 82%.




Figure 6: Representative position traces, one episode per cell of the main ablation. For each cell, the episode shown is the one whose survival is closest to the cell’s median across all 400 evaluation episodes (4 training seeds $\times$ 100 evaluation worlds). The path is colored by tick (light early, dark late). The start position is the open circle. Food-eaten events along the path are green dots. The end marker indicates cause of death (red “X” for enemy contact, gray “X” for starvation, green square for reaching the time cap). Food and shelter locations are not drawn, since food respawns stochastically and shelters are unoccupied in every DQN episode.

For the MINIMAL configuration this is informative, because the reward was specifically designed to teach the heuristic’s shelter mode via a shelter-proximity reward gated on hunger ratio $\geq u_{\text{gate}}$ (Section 6.3). One plausible hypothesis is that the gate never opens, in which case the agent’s hunger ratio stays below $u_{\text{gate}} = 0.5$ for most of the episode and the shelter-proximity reward never fires. To test this, we logged the agent’s hunger ratio at every tick and ran the four trained MINIMAL_FS checkpoints on the 100-world evaluation suite. Across these 400 episodes the agent’s hunger ratio was at or above u_{gate} for a mean of 32.0% of episode ticks (longest continuous open-gate streak averaged 237 ticks per episode, median 249, max 324), and all 400 episodes started above the gate since the agent begins every episode at full hunger. The gate is therefore open often enough and for long enough that the shelter-proximity reward has substantial opportunity to fire, and this hypothesis is rejected.

An alternative explanation is that the $\Delta\phi$ formulation of the shelter-proximity reward pays only for approaching a shelter, not for occupying one. Once the agent is inside a shelter the closeness no longer changes, so $\Delta\phi^{(s)} = 0$ on every subsequent tick, and the shelter-safety reward w_{shelter} , which pays for occupancy directly, is zero in the present

MINIMAL configuration (Table 3). On its own this cannot explain the gap, because the four misspecified cells keep w_{shelter} at its default +0.1 and show the same near-zero occupancy. What no cell provides is both halves at once. The misspecified rewards pay for being sheltered but give no directional signal for finding a shelter, so the agent could only discover the payment through undirected exploration, and the payment itself fires only when an enemy is within distance 3 of the sheltered agent (Table 1), an event a passing shelter visit rarely triggers. MINIMAL on the other hand gives the directional signal, but pays nothing for staying.

Finally, shelter use may not even be strictly necessary for survival in this environment, since the enemy density and pursuit rules of Section 4.1 may simply not impose enough damage pressure to make shelter-camping a required strategy. The fact that MINIMAL_FS reaches 762 ticks, similar to the heuristic, without any shelter use supports this possibility.

Beyond the shelter question, the BASELINE versus ABSOLUTE_PROXIMITY contrast isolates the effect of the closeness-change reward shape. At single-frame, switching the proximity term from $\Delta\phi$ to ϕ degrades survival by 13 ticks ($647 \rightarrow 634$) and lowers food intake from 0.66 to 0.43 per episode. The lower food intake is consistent with the agent earning proximity reward for being near food without crossing onto the food tile to eat it. At frame-stacked, the same switch improves survival by 32 ticks ($555 \rightarrow 587$) with essentially unchanged food intake ($4.83 \rightarrow 4.82$). The 4-frame stack appears to give the agent enough temporal information to navigate onto food rather than hover next to it, so the absolute form is no worse than the closeness-change form at this observation level. The closeness-change form, used in MINIMAL, remains the only reward shape under which frame stacking improves survival relative to its own single-frame counterpart.

7.3 Learning dynamics

Figure 7 plots the in-training 100-episode evaluation survival as a function of training step for each of the six ablation cells. The trajectories support three observations relevant to the experimental design. First, each cell’s curve flattens by step 1.0–1.4M, and the second half of training contributes drift rather than a discernible upward trend in any cell. Because of this, we terminate training at $T = 2,000,000$ steps. We cannot rule out “grokking”, a sudden jump in survival after long stagnation [Power et al., 2022] or a slow but steady improvement past this budget. Second, the trajectories are tightly clustered across seeds within a cell, which corroborates the per-seed survival homogeneity reported in Table 7 and indicates that the variance across cells is attributable to (reward, observation) variation rather than random initialization. Third, the displayed values use a 10-eval trailing mean (partial window for the first 9 evals) over 100 episodes per eval, so they track the 100-episode benchmark numbers reported in Table 4 closely.

7.4 Limitations

The most important limitation is external validity. All findings come from a single, custom-built environment and a single task, so the numbers are not directly comparable to results reported for Crafter, MiniGrid, or other gridworld libraries, and the reward-dependent direction of the frame-stacking effect can only be generalized through replication on other environments.

The second limitation concerns the design of the reward axis itself. MINIMAL differs from BASELINE in several ways at once: it removes the hunger penalties, drops the event




Figure 7: In-training evaluation survival vs training step for the six ablation cells. Each line is one seed, smoothed with a 10-eval rolling mean (50,000-step window). The dashed gray line at 769 ticks marks the heuristic baseline.

rewards for eating and taking damage while keeping only the death penalty, activates proximity terms that BASELINE leaves disabled, gates two of them on hunger, and assigns different weights. “Clean versus leaky” is therefore one description of a many-dimensional change, not a single isolated factor. The return signature of Section 7.1 supports the leak-exploitation mechanism, since under frame stacking the leaky cells collect more reward while surviving shorter, but the design cannot rule out that some other ingredient of the bundle is the actual causal variable. An ablation that removed one leak at a time, for example zeroing the hunger-proportional and starvation penalties, the pair that makes a quick death cheaper than slow starvation, while leaving the rest of BASELINE intact, would isolate the mechanism more sharply.

Three further constraints qualify the strength of the claims. First, the network architecture is held fixed across observation depths, so we do not run an architecture sweep. The reward evidence in Section 7.1 rules out an undersized $k = 4$ network as the driver of the leaky-cell degradation (the effect flips sign with reward shape, and the leaky cells gain reward rather than lose it, neither of which an under-provisioned agent would do),

but training the $k = 4$ cells with a network scaled to their larger input would confirm this more directly. Second, the computational budget limited each cell to two million training steps. The learning curves flatten well before this point, but longer training could test whether certain configurations improve further. Third, each cell uses four training seeds, below the five to ten typically recommended in RL methodology work [Henderson et al., 2018, Colas et al., 2018]. The within-cell standard errors in Table 4 are computed from this small sample and therefore carry substantial estimation noise themselves. A larger seed budget would estimate each cell mean more precisely and show more directly whether the survival differences between the single-frame and frame-stacked variants hold up when training is repeated with new random seeds. Finally, WEAK_PROXIMITY was trained on a single seed, so we draw no conclusions from it beyond the qualitative note in Section 6.4.

8 Conclusion

We ran a $3 \times 2 \times 4$ ablation that crossed three reward configurations (BASELINE, ABSOLUTE_PROXIMITY, MINIMAL) with two observation types (single-frame and 4-frame stacked), four random seeds per cell. We added a single-seed illustrative run of WEAK_PROXIMITY to demonstrate the suicide failure mode under a too-small proximity weight.

We established three predictions: that MINIMAL_FS would be the best-performing cell (P1), that it would exceed the heuristic (P2), and that frame stacking would improve every reward configuration over its single-frame counterpart (P3). P1 held. MINIMAL_FS reaches mean survival 762 ticks, the highest of any cell. P2 is not rejected but also not affirmed. The paired gap to the heuristic is -10.0 ticks (95% CI $[-36, +16]$, $p = 0.45$), so MINIMAL_FS could be tied with the heuristic but we cannot show it exceeds the heuristic. P3 did not hold: while frame stacking improves MINIMAL (+78 ticks), it degrades ABSOLUTE_PROXIMITY (-47) and BASELINE (-92).

Beyond refuting P3, the data reveal a directional structure to frame stacking’s effect. It improves the cleanest reward configuration and degrades both leakier ones, by margins much larger than within-cell variance. Although we did not set out to test this, it emerged from the ablation, and is the most directly informative result of the work. It showed that adding observation history is not an unconditional improvement on this task, and that under the two leakier reward designs studied here, it even actively hurt the resulting agent. We discussed plausible mechanisms in Section 7.

This work studies the reward $\times$ observation interaction in DQN on one specific environment. The main contribution is the $3 \times 2 \times 4$ ablation just described, which tests three pre-specified predictions and reveals that the sign of the frame-stacking effect depends on reward shape. A second finding confirms that for the two rewards that make a quick death cheaper than slow starvation, 19–25% of evaluation episodes end with the agent closing distance on the enemy that kills it, while MINIMAL, which equalizes the reward cost of fast and slow death, shows at most 2%.

As a secondary outcome, the leaky-reward results corroborate the reward-capability interaction of Pan et al. [2022] in a new setting: the same direction they document across model size, action-space resolution, observation noise, and training time also appears when capability is varied through observation history, in a different environment. We do not, however, resolve whether the survival drop has the phase-transition structure they report, since our design varies k at only two levels.

Four directions follow naturally from these results. First, replicate the ablation on

a second multi-objective environment (for example a forked MiniGrid or Crafter setup) to test the generality of the reward-dependent direction. Second, train the $k = 4$ cells with a network scaled up to match the larger input. The current architecture is the same for $k = 1$ and $k = 4$, despite the $k = 4$ observation being roughly four times the size of the $k = 1$ observation. Scaling up the $k = 4$ DQN agent to match the input size would help confirm that reward design, not missing capacity, explains the decrease in survival. Third, train all cells for substantially longer than two million steps to test for further survival improvements. Fourth, test whether adding a shelter occupancy reward to MINIMAL produces learned shelter use. MINIMAL would then be the first cell to combine a directional signal for finding shelters with a payment for staying in one, the combination Section 7.2 identifies as untested.

A Hyperparameters

Table 6 lists all agent and training hyperparameters used across the ablation. Values follow established defaults from the literature rather than environment-specific tuning. The frame-stack size k is the only varying parameter, all others are held constant across all twenty-four ablation runs.

Table 6: Agent hyperparameters.

Parameter	Value	Reference
<i>Optimization</i>		
Learning rate	10^{-4}	after Loshchilov and Hutter [2019]
Weight decay	10^{-4}	
Batch size	128	
Gradient clip norm	1.0	
<i>Replay buffer</i>		
Buffer capacity	250,000	after Mnih et al. [2015]
Min. buffer size (warmup)	10,000	
PER α	0.6	
PER β_0	0.4	Schaul et al. [2016]
β annealing steps	500,000	Hessel et al. [2018]
n -step returns	3	
<i>Target network</i>		
Polyak τ	0.002	after Lillicrap et al. [2016]
<i>Exploration</i>		
ε start / end	1.0 / 0.01	after Mnih et al. [2015]
ε decay steps	500,000	
<i>Network</i>		
CNN channels	(32, 64, 64)	Mnih et al. [2015]
CNN kernel sizes	(3×3, 3×3, 3×3)	
CNN strides	(1, 1, 2)	
CNN padding	(1, 1, 0)	Ba et al. [2016]
Normalization	LayerNorm	
Merge hidden dim	256	
Dueling head hidden dim	128	
<i>Observation (ablation variable)</i>		
Frame stack k	{1, 4}	Mnih et al. [2015]
<i>Training</i>		
Discount γ	0.99	
Total timesteps T	2,000,000	
Eval frequency	5,000 steps	
Eval episodes (in-training)	100	
Final benchmark episodes	100 (seeds 1000–1099)	

The convolutional encoder reduces the 15×15 spatial grid to 7×7 at the third layer, producing a 3136-dimensional feature vector. This is concatenated with the three scalar observations and projected through the merge layer to a 256-dimensional representation. LayerNorm is applied after the first two convolutions and after the merge layer. The Dueling

head splits this representation into a value stream ($256 \rightarrow 128 \rightarrow 1$) and an advantage stream ($256 \rightarrow 128 \rightarrow 5$), and Q -values are computed as $Q(s, a) = V(s) + A(s, a) - \bar{A}(s)$. For the 4-frame stacked cells ($k = 4$), the only structural change is the input depth of the first convolution ($12 \rightarrow 32$ instead of $3 \rightarrow 32$).

B Per-seed survival

Table 7: Per-seed DQN survival (mean over the 100-episode benchmark per seed). Heuristic survival is 769 ± 19 ticks. Cell-level uncertainty and the comparison to the heuristic are in Table 4.

Cell	seed 42	seed 43	seed 44	seed 45
BASELINE	623	698	631	636
BASELINE_FS	561	526	546	586
ABSOLUTE_PROXIMITY	625	651	645	616
ABSOLUTE_PROXIMITY_FS	564	587	595	600
MINIMAL	684	681	678	694
MINIMAL_FS	755	760	790	744

C Episode outcomes and shelter occupancy

Table 8: Per-cell episode outcomes and shelter occupancy, in percent. Outcomes are mutually exclusive and sum to 100% within rounding. “Sh. occ.” is the mean fraction of survived ticks spent inside a shelter, expressed as a percentage. Each row averages four seeds (the heuristic row is a single 100-episode run).

Cell	Time cap	Enemy suicide	Failure to escape	Starvation	Sh. occ.
Heuristic	6.0	2.0	0.0	92.0	82.1
<i>Single-frame ($k = 1$)</i>					
BASELINE	1.5	22.5	11.5	64.5	0.9
ABSOLUTE_PROXIMITY	0.5	19.5	14.5	65.5	1.0
MINIMAL	0.0	0.0	6.8	93.2	0.6
<i>4-frame stacked ($k = 4$)</i>					
BASELINE_FS	2.8	24.8	17.8	54.8	0.6
ABSOLUTE_PROXIMITY_FS	4.2	18.5	14.8	62.5	1.4
MINIMAL_FS	0.0	1.8	2.0	96.2	1.0

References

- Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete problems in ai safety. *arXiv preprint arXiv:1606.06565*, 2016. doi: 10.48550/arXiv.1606.06565.
- Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E. Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.
- Maxime Chevalier-Boisvert, Bolun Dai, Mark Towers, Rodrigo de Lazcano, Lucas Willems, Salem Lahlou, Suman Pal, Pablo Samuel Castro, and Jordan Terry. Minigrid & miniworld: Modular & customizable reinforcement learning environments for goal-oriented tasks. In *Advances in Neural Information Processing Systems (Datasets and Benchmarks Track)*, 2023.
- Tyler Clark, Mark Towers, Christine Evers, and Jonathon Hare. Beyond the rainbow: High performance deep reinforcement learning on a desktop PC. In *Proceedings of the International Conference on Machine Learning*, 2025.
- Cédric Colas, Olivier Sigaud, and Pierre-Yves Oudeyer. How many random seeds? Statistical power analysis in deep reinforcement learning experiments. *arXiv preprint arXiv:1806.08295*, 2018.
- Matteo Gallici, Mattie Fellows, Benjamin Ellis, Bartomeu Pou, Ivan Masmitja, Jakob Nicolaus Foerster, and Mario Martin. Simplifying deep temporal difference learning. *arXiv preprint arXiv:2407.04811*, 2024.
- Danijar Hafner. Benchmarking the spectrum of agent capabilities. In *Proceedings of the International Conference on Learning Representations*, 2022.
- Charles R. Harris, K. Jarrod Millman, Stéfan J. van der Walt, Ralf Gommers, Pauli Virtanen, David Cournapeau, Eric Wieser, Julian Taylor, Sebastian Berg, Nathaniel J. Smith, Robert Kern, Matti Picus, Stephan Hoyer, Marten H. van Kerkwijk, Matthew Brett, Allan Haldane, Jaime Fernández del Río, Mark Wiebe, Pearu Peterson, Pierre Gérard-Marchant, Kevin Sheppard, Tyler Reddy, Warren Weckesser, Hameer Abbasi, Christoph Gohlke, and Travis E. Oliphant. Array programming with NumPy. *Nature*, 585(7825):357–362, 2020. doi: 10.1038/s41586-020-2649-2.
- Matthew Hausknecht and Peter Stone. Deep recurrent Q-learning for partially observable MDPs. In *AAAI Fall Symposium on Sequential Decision Making for Intelligent Agents*, 2015.
- Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger. Deep reinforcement learning that matters. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 32, 2018.
- Matteo Hessel, Joseph Modayil, Hado van Hasselt, Tom Schaul, Georg Ostrovski, Will Dabney, Dan Horgan, Bilal Piot, Mohammad Azar, and David Silver. Rainbow: Combining improvements in deep reinforcement learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2018.

- Jemin Hwangbo, Inkyu Sa, Roland Siegwart, and Marco Hutter. Control of a quadrotor with reinforcement learning. *IEEE Robotics and Automation Letters*, 2(4):2096–2103, 2017.
- Jemin Hwangbo, Joonho Lee, Alexey Dosovitskiy, Dario Bellicoso, Vassilios Tsounis, Vladlen Koltun, and Marco Hutter. Learning agile and dynamic motor skills for legged robots. *Science Robotics*, 4(26), 2019.
- Heinrich Küttler, Nantas Nardelli, Alexander H. Miller, Roberta Raileanu, Marco Selvatici, Edward Grefenstette, and Tim Rocktäschel. The NetHack learning environment. In *Advances in Neural Information Processing Systems*, 2020.
- Joel Lehman, Jeff Clune, Dusan Misevic, Christoph Adami, Lee Altenberg, Julie Beaulieu, Peter J. Bentley, Samuel Bernard, Guillaume Beslon, David M. Bryson, Nick Cheney, Patryk Chrabaszcz, Antoine Cully, Stephane Doncieux, Fred C. Dyer, Kai Olav Ellefsen, Robert Feldt, Stephan Fischer, Stephanie Forrest, Antoine Frénoy, Christian Gagné, Leni Le Goff, Laura M. Grabowski, Babak Hodjat, Frank Hutter, Laurent Keller, Carole Knibbe, Peter Krcak, Richard E. Lenski, Hod Lipson, Robert MacCurdy, Carlos Maestre, Risto Miikkulainen, Sara Mitri, David E. Moriarty, Jean-Baptiste Mouret, Anh Nguyen, Charles Ofria, Marc Parizeau, David Parsons, Robert T. Pennock, William F. Punch, Thomas S. Ray, Marc Schoenauer, Eric Schulte, Karl Sims, Kenneth O. Stanley, François Taddei, Danesh Tarapore, Simon Thibault, Richard Watson, Westley Weimer, and Jason Yosinski. The surprising creativity of digital evolution: A collection of anecdotes from the evolutionary computation and artificial life research communities. *Artificial Life*, 26(2):274–306, 2020. doi: 10.1162/artl_a_00319.
- Jan Leike, Miljan Martic, Victoria Krakovna, Pedro A. Ortega, Tom Everitt, Andrew Lefrancq, Laurent Orseau, and Shane Legg. AI safety gridworlds. *arXiv preprint arXiv:1711.09883*, 2017.
- Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. In *Proceedings of the International Conference on Learning Representations*, 2016.
- Antonio Loquercio, Elia Kaufmann, René Ranftl, Matthias Müller, Vladlen Koltun, and Davide Scaramuzza. Learning high-speed flight in the wild. *Science Robotics*, 6(59), 2021.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *Proceedings of the International Conference on Learning Representations*, 2019.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin Riedmiller. Playing Atari with deep reinforcement learning. *arXiv preprint arXiv:1312.5602*, 2013.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540): 529–533, 2015.

- Johan Obando-Ceron and Pablo Samuel Castro. Revisiting rainbow: Promoting more insightful and inclusive deep RL research. In *Proceedings of the International Conference on Machine Learning*, 2021.
- Johan Obando-Ceron, Marc G. Bellemare, and Pablo Samuel Castro. Small batch deep reinforcement learning. In *Advances in Neural Information Processing Systems*, 2023.
- Alexander Pan, Kush Bhatia, and Jacob Steinhardt. The effects of reward misspecification: Mapping and mitigating misaligned models. In *Proceedings of the International Conference on Learning Representations*, 2022.
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems*, volume 32, 2019.
- Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*, 2022.
- Tom Schaul, John Quan, Ioannis Antonoglou, and David Silver. Prioritized experience replay. In *Proceedings of the International Conference on Learning Representations*, 2016.
- Max Schwarzer, Johan Obando-Ceron, Aaron Courville, Marc G. Bellemare, Rishabh Agarwal, and Pablo Samuel Castro. Bigger, better, faster: Human-level Atari with human-level efficiency. In *Proceedings of the International Conference on Machine Learning*, 2023.
- Richard S. Sutton. Learning to predict by the methods of temporal differences. *Machine Learning*, 3(1):9–44, 1988.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018.
- Hado van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double Q-learning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2016.
- Ziyu Wang, Tom Schaul, Matteo Hessel, Hado van Hasselt, Marc Lanctot, and Nando de Freitas. Dueling network architectures for deep reinforcement learning. In *Proceedings of the International Conference on Machine Learning*, 2016.